

# Experiment 8 - Introduction to Computer Vision

Computer vision is a multidisciplinary field that enables machines to interpret and understand visual data, such as images and videos. In this experiment, we explore different fields of computer vision, starting with basics of image manipulation, and going to more advanced problems such as object recognition and detection. We will mainly use OpenCV, Scikit-learn, and PyTorch Python packages for the procedure of this experiment. Parts of this experiment is based on PyTorch tutorials.

## 1.1 Image processing and manipulation with OpenCV

In this part, we will gain a foundational understanding of image processing and manipulation using the OpenCV library in Python. To install the OpenCV library, you can use the following command

---

```
$ pip install opencv-python
```

---

In the following sections, we will use the Birzeit University campus picture shown in Figure 1.1 as an example. To download the picture you can use the command

---

```
$ wget https://www.birzeit.edu/sites/default/files/buildings.jpg
```

---

Kindly note that if you are running the commands above in google colab, then you need to add ! at the beginning of the command.



Figure 1.1: BZU Campus.

## Loading and displaying an image

To load the image we downloaded to numpy array and display it, use the following code

---

Code Snippet 1.1: Loading and displaying an image.

---

```
import cv2 as cv
import matplotlib.pyplot as plt

# Load an image
image = cv.imread('buildings.jpg')

# Display the image
plt.imshow(cv.cvtColor(image, cv.COLOR_BGR2RGB))
plt.title('Original Image')
plt.show()
```

---

**Task 1:** What is the shape of the loaded image? What does each dimension represent?

**Task 2:** repeat the previous code but now pass the image directly to `plt.imshow()`. I.e., without using `cv.cvtColor()`. Do you notice any difference? Why?

## Grayscale conversion

To convert the image in the previous section into grayscale, use the code below

---

Code Snippet 1.2: Grayscale conversion.

---

```
# Convert the image to grayscale
gray_image = cv.cvtColor(image, cv.COLOR_BGR2GRAY)

# Display the grayscale image
plt.imshow(gray_image, cmap='gray')
plt.title('Grayscale Image')
plt.show()
```

---

**Task 3:** What is the shape of the `gray_image`? What does each dimension represent?

## Image transformation

Run the code below which applies some transformations on the image. Try to explain each of the applied transformations

---

Code Snippet 1.3: Image transformation.

---

```
image = cv.cvtColor(image, cv.COLOR_BGR2RGB)
# Resize and crop the image
resized_image = cv.resize(image, (300, 200))
cropped_image = image[50:250, 50:250]

# Rotate the image
rows, cols = image.shape[:2]
```

```

M = cv.getRotationMatrix2D((cols / 2, rows / 2), 45, 1)
rotated_image = cv.warpAffine(image, M, (cols, rows))

# Display the transformed images
plt.imshow(resized_image)
plt.title('Resized Image')
plt.show()

plt.imshow(cropped_image)
plt.title('Cropped Image')
plt.show()

plt.imshow(rotated_image)
plt.title('Rotated Image')
plt.show()

```

---

**Task 4:** Rotate the original image by 10 degrees counter-clockwise around the upper-left corner of the image.

### Gaussian blurring

Gaussian blurring is used to reduce the noise in the image. However, it also affects the details of the input image. The basic idea is to convolve the image with a Gaussian kernel, which is a 2D bell-shaped function. The following code can be used to apply Gaussian blurring

---

Code Snippet 1.4: Gaussian blurring.

---

```

# Apply Gaussian blur
blurred_image = cv.GaussianBlur(image, (7, 7), 0)

# Display the result
plt.imshow(blurred_image)
plt.title('Blurred Image')
plt.show()

plt.imshow(image)
plt.title('Original Image')
plt.show()

```

---

**Task 5:** What does (7, 7) in the previous code represent? try to increase this value, what do you notice?

### Edge detection

To detect edges in the image, run the following code

---

Code Snippet 1.5: Edge detection.

---

```

# Perform edge detection using Canny
edges = cv.Canny(gray_image, 50, 150)

```

```
# Display the result
plt.imshow(edges, cmap='gray')
plt.title('Edge Detection')
plt.show()
```

---

**Task 6:** Explain the parameters of `cv.Canny()`. How do they affect the results?

## 1.2 Image Classification / Object Recognition

Image classification is a fundamental task in computer vision. The goal of image classification is to assign a single label for an input image. Typically, the image contains only one object, and the goal is to recognize which object is depicted in the image.

For this lab, we will use the CIFAR10 dataset, which is a standard benchmark for image classification. The dataset consists of 10 classes: ‘airplane’, ‘automobile’, ‘bird’, ‘cat’, ‘deer’, ‘dog’, ‘frog’, ‘horse’, ‘ship’, ‘truck’. Each image in CIFAR10 are of size  $3 \times 32 \times 32$ , i.e. 3-channel color images of size  $32 \times 32$  pixels.

### 1.2.1 Image classification with MLP

In this section, we will train an MLP classifier on CIFAR10 dataset. Let’s start by loading the dataset. Luckily, PyTorch has a package called `torchvision`, that has data loaders for common datasets including CIFAR10, and data transformers for images `torchvision.datasets` and `torch.utils.data.DataLoader`.

#### Loading images

The first step is to load and normalize the CIFAR10 training and test images using `torchvision`.

---

Code Snippet 1.6: Loading and normalizing CIFAR10 images.

---

```
import torch
import torchvision
import torchvision.transforms as transforms

transform = transforms.Compose(
    [transforms.ToTensor(),
     transforms.Normalize((0.5, 0.5, 0.5), (0.5, 0.5, 0.5))])

batch_size = 4

trainset = torchvision.datasets.CIFAR10(root='./data', train=True,
                                         download=True, transform=transform)
trainloader = torch.utils.data.DataLoader(trainset, batch_size=batch_size,
```

```

shuffle=True, num_workers=2)

testset = torchvision.datasets.CIFAR10(root='./data', train=False,
                                       download=True, transform=transform)
testloader = torch.utils.data.DataLoader(testset, batch_size=batch_size,
                                       shuffle=False, num_workers=2)

classes = ('plane', 'car', 'bird', 'cat',
           'deer', 'dog', 'frog', 'horse', 'ship', 'truck')

```

---

The output of torchvision datasets are PILImage images of range [0, 1]. We transform them to Tensors of normalized range [-1, 1]. The following code shows some of the training images.

Code Snippet 1.7: Displaying some training images.

---

```

import matplotlib.pyplot as plt
import numpy as np

# functions to show an image

def imshow(img):
    img = img / 2 + 0.5 # unnormalize
    npimg = img.numpy()
    plt.imshow(np.transpose(npimg, (1, 2, 0)))
    plt.show()

# get some random training images
dataiter = iter(trainloader)
images, labels = next(dataiter)

# show images
imshow(torchvision.utils.make_grid(images))
# print labels
print(' '.join(f'{classes[labels[j]]:5s}' for j in range(batch_size)))

```

---

**Task 7:** Explain the effect of the following transforms `transforms.ToTensor()`, and `transforms.Normalize((0.5, 0.5, 0.5), (0.5, 0.5, 0.5))`.

**Task 8:** What is the purpose of the following line: `np.transpose(npimg, (1, 2, 0))`?

## Defining the network

Let's define a simple MLP. This network is the same as the MLP from experiment 7, except for the input layer, which has more inputs.

---

```

import torch
import torch.nn as nn

```

```

import torch.nn.functional as F

class Net(nn.Module):

    def __init__(self):
        super(Net, self).__init__()
        self.fc1 = nn.Linear(3072, 64)
        self.fc2 = nn.Linear(64, 64)
        self.fc3 = nn.Linear(64, 10)

    def forward(self, x):
        # apply the first layer with relu activation
        x = F.relu(self.fc1(x))
        x = F.relu(self.fc2(x))
        x = self.fc3(x)
        return x

net = Net()

device = torch.device('cuda:0' if torch.cuda.is_available() else 'cpu')
# Assuming that we are on a CUDA machine, this should print a CUDA device:
print(device)

net.to(device)

```

---

## Define a Loss function and optimizer

Let's use a Classification Cross-Entropy loss and SGD with momentum.

---

```

import torch.optim as optim

criterion = nn.CrossEntropyLoss()
optimizer = optim.SGD(net.parameters(), lr=0.001, momentum=0.9)

```

---

## Train the network

---

```

net.train()
for epoch in range(2): # loop over the dataset multiple times

    running_loss = 0.0
    for i, data in enumerate(trainloader, 0):
        # get the inputs; data is a list of [inputs, labels]
        inputs, labels = data
        inputs = inputs.to(device)
        labels = labels.to(device)

```

```

# zero the parameter gradients
optimizer.zero_grad()

# forward + backward + optimize
outputs = net(torch.flatten(inputs,1))
loss = criterion(outputs, labels)
loss.backward()
optimizer.step()

# print statistics
running_loss += loss.item()
if i % 2000 == 1999: # print every 2000 mini-batches
    print(f'[{epoch + 1}, {i + 1:5d}] loss: {running_loss /
        2000:.3f}')
    running_loss = 0.0

print('Finished Training')

```

---

### Test the network on the test data

We have trained the network for 2 passes over the training dataset. But we need to check if the network has learnt anything at all.

We will check this by predicting the class label that the neural network outputs, and checking it against the ground-truth. If the prediction is correct, we add the sample to the list of correct predictions.

---

```

correct = 0
total = 0
# since we're not training, we don't need to calculate the gradients for
# our outputs
net.eval()
with torch.no_grad():
    for data in testloader:
        images, labels = data
        images, labels = images.to(device), labels.to(device)
        # calculate outputs by running images through the network
        outputs = net(torch.flatten(images,1))
        # the class with the highest energy is what we choose as prediction
        _, predicted = torch.max(outputs.data, 1)
        total += labels.size(0)
        correct += (predicted == labels).sum().item()

print(f'Accuracy of the network on the 10000 test images: {100 * correct
    // total} %')

```

---

**Task 9:** For the first layer in the MLP (`self.fc1 = nn.Linear(3072, 64)`), what does the number 3072 represent?

### 1.2.2 Image classification with CNNs

In the previous section, we used a MLP to classify images. However, for this type of data, convolutional neural networks (CNNs) are a better choice. Convolutional Neural Networks (CNNs) are a class of deep learning models designed specifically for processing structured grid data, such as images. The key innovation of CNNs lies in their ability to automatically and adaptively learn hierarchical representations of input data. The fundamental building blocks of CNNs are convolutional layers, which apply convolution operations to the input data. Convolution involves sliding a small filter (also called a kernel) across the input, performing element-wise multiplications, and aggregating the results to create a feature map.

CNNs excel at capturing spatial hierarchies and local patterns within the input data. Convolutional layers are typically followed by activation functions (e.g., ReLU) to introduce non-linearity and pooling layers to downsample the spatial dimensions, reducing the computational load and preserving important features. The final layers of a CNN typically include fully connected layers to make predictions based on the learned hierarchical representations. CNNs have proven highly effective in various computer vision tasks, such as image classification, object detection, and image segmentation, due to their ability to automatically learn relevant features from raw pixel data. Figure 1.2 shows the general architecture of CNNs.

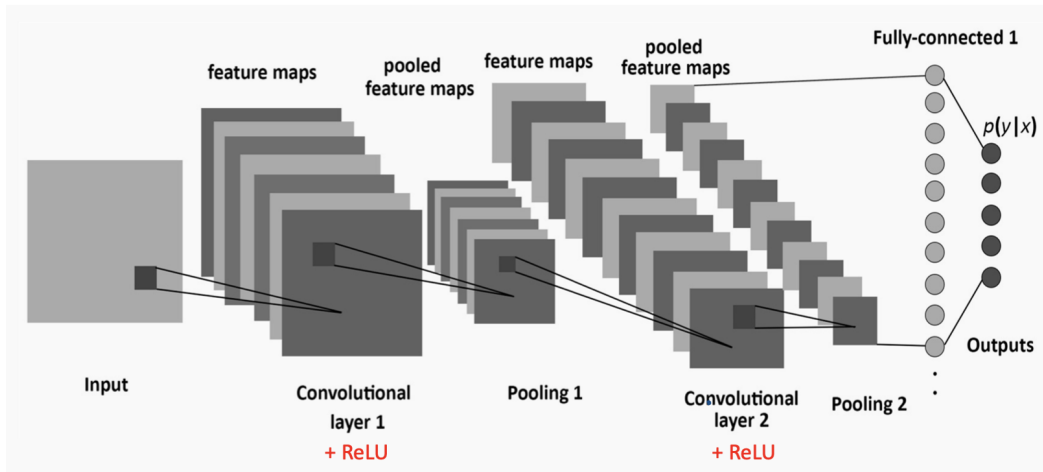


Figure 1.2: General architecture of convolutional neural networks.

As shown in the figure, CNNs usually contains three types of layers:

- **Convolutional Layer:** The convolutional layer is the core building block of a CNN. It applies convolution operations to the input data using filters or kernels. These filters slide over the input, performing element-wise multiplications and aggregating the results to create feature maps. This process allows the network to automatically learn spatial hierarchies and local patterns, capturing relevant features in the input data. Figure 1.3 shows a convolutional layer with 2 filters.
- **Pooling Layer:** Pooling layers are used to downsample the spatial dimensions of the feature maps generated by the convolutional layers. Max pooling and






Figure 1.3: A convolutional layer with 2 filters.

average pooling are common techniques, where the operation involves taking the maximum or average value from a group of neighboring pixels. Pooling helps reduce the computational complexity of the network, makes the learned features more invariant to small translations, and retains the most important information. Figure 1.4 shows both the max and average pooling operations.

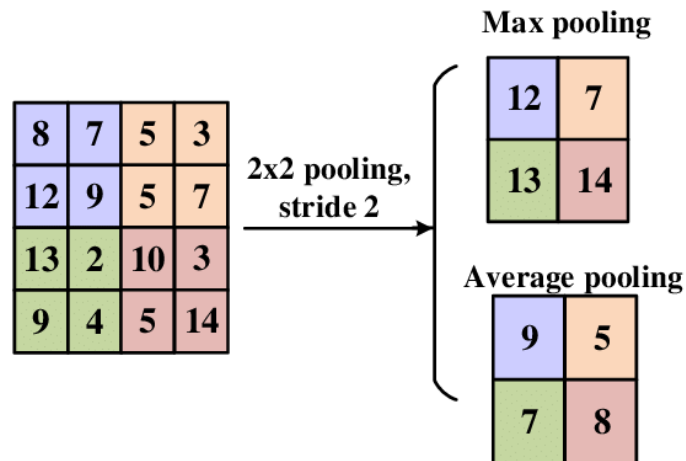


Figure 1.4: The operation of max and average pooling.

- **Fully Connected Layer:** Fully connected layers, also known as dense layers, are traditionally found at the end of the CNN architecture. Neurons in a fully connected layer connect to all the neurons in the previous layer, effectively creating a dense matrix of connections. These layers are responsible for combining the high-level features learned by the convolutional layers to make final predictions or classifications. In image classification tasks, the output of the last fully connected layer is often fed into a softmax activation function to produce probability scores for different classes.

Let's now train a simple CNN to classify images from CIFAR10 dataset. Repeat the procedure from the previous section but replace the network with the CNN defined in the following code snippet. Note that CNNs expects a 3-dimensional tensor at the input, which means that we have to pass the input image without flattening both during training and testing. Make sure that you adjust the code from the previous section accordingly.

---

```
import torch.nn as nn
import torch.nn.functional as F

class NetCNN(nn.Module):
    def __init__(self):
        super().__init__()
        self.conv1 = nn.Conv2d(3, 6, 5)
        self.pool = nn.MaxPool2d(2, 2)
        self.conv2 = nn.Conv2d(6, 16, 5)
        self.fc1 = nn.Linear(16 * 5 * 5, 120)
        self.fc2 = nn.Linear(120, 84)
        self.fc3 = nn.Linear(84, 10)

    def forward(self, x):
        x = self.pool(F.relu(self.conv1(x)))
        x = self.pool(F.relu(self.conv2(x)))
        x = torch.flatten(x, 1) # flatten all dimensions except batch
        x = F.relu(self.fc1(x))
        x = F.relu(self.fc2(x))
        x = self.fc3(x)
        return x

net = NetCNN()
net = net.to(device)
```

---

**Task 10:** Study the CNN code and answer the following questions:

- How many layers does the network have?
- What does the parameters of `nn.Conv2d` represent?
- What does the parameters of `nn.MaxPool2d` represent?
- In the first linear layer, what does the number  $16 * 5 * 5$  represent? How to obtain this number?

### 1.2.3 Transfer Learning

In previous sections, we worked with simple networks, but in practical applications, more extensive networks with millions of parameters are often employed. Training

such complex models from scratch on small datasets tends to lead to overfitting. To avoid this problem, transfer learning can be used. In the transfer learning paradigm, a model pre-trained on a task is utilized to improve performance on a different yet related task. The approach involves training the model on a task where a large dataset is available and subsequently applying this acquired knowledge to a new, often smaller, dataset. There are two main approaches for transfer learning:

- Finetuning: Instead of random initialization, we initialize the network with a pretrained network. Rest of the training looks as usual.
- ConvNet as fixed feature extractor: Here, we freeze the weights for all of the network except that of the final fully connected layer. This last fully connected layer is replaced with a new one with random weights and only this layer is trained.

### Finetuning the ConvNet

In this section, we will finetune AlexNet, which is a commonly used network for image related tasks. To this end, repeat the training and testing procedure from the previous section to train AlexNet model pre-trained on the imagenet dataset. To define the network, you can use the following code. Make sure to use GPU for this part, otherwise, training would take very long time.

---

```
from torchvision import datasets, models

net = models.alexnet(weights='IMAGENET1K_V1')
net.classifier[6] = nn.Linear(4096, 10)
net = net.to(device)
```

---

Make sure also to apply the following transformation on the data

---

```
transform = transforms.Compose(
    [transforms.Resize(256),
     transforms.CenterCrop(224),
     transforms.ToTensor(),
     transforms.Normalize([0.485, 0.456, 0.406], [0.229, 0.224, 0.225])])
```

---

### Finetuning vs Learning from scratch

Repeat the previous part but now without using the pretrained weights. What do you notice?

To use AlexNet with random initialization you can use the following code to define the model

---

```
from torchvision import datasets, models

net = models.alexnet()
net.classifier[6] = nn.Linear(4096, 10)
net = net.to(device)
```

---

## 1.3 Case Study: Object Detection

Object detection is a computer vision task that involves identifying and locating objects within an image or a video sequence. Unlike image classification, which assigns a single label to an entire image, object detection aims to recognize and localize multiple objects, each with its corresponding bounding box, in the given input. Figure 1.5 shows a sample output for an object detector.

In this case study, we will study and compare different object detection models. The details and requirements for the task will be given by the instructor.



Figure 1.5: The object detection task.